

Assignment #3 (60pts + 25 bonus points)

Document Classification using Naïve Bayes

Out: 5/5, 2020

Due: 5/15-5/17 midnight (11:59pm).

Submission

1. All files should be submitted to Canvas. You can submit multiple times, but please have the same team member resubmit all required files each time.
2. Implement in Python.
3. Submit the **assn3.zip** file. Submit your programs (if any) and a report (**report.**txt, .doc, or .pdf) that describes all your work including all the assumptions you made, your output/results, and the discussion of the results if possible.
4. **DO NOT** include the given data files in your submission!!
5. **The submission link will be closed at 11:59pm 5/17 (Friday). No late submission is accepted.**

In this assignment, you will be using Naïve Bayes to classify email messages into spam and nonspam groups.

Step 1: Familiarize the Data Set

The data set is a preprocessed subset of the Ling-Spam Dataset that is based on 960 real email messages from a linguistics mailing list.

The dataset is split into two subsets: a 700-email subset for training and a 260-email subset for testing. Each of the training and testing subsets contain 50% spam messages and 50% nonspam messages. Additionally, the emails have been preprocessed in the following ways:

1. Stop word removal: Certain words like "and," "the," and "of," are very common in all English sentences and are not very meaningful in deciding spam/nonspam status, so these words have been removed from the emails.
2. Lemmatization: Words that have the same meaning but different endings have been adjusted so that they all have the same form. For example, "include", "includes," and "included," would all be represented as "include." All words in the email body have also been converted to lower case.
3. Removal of non-words: Numbers and punctuation have both been removed. All white spaces (tabs, newlines, spaces) have all been trimmed to a single space character.

As you can discover from browsing these messages, preprocessing has left occasional word fragments and nonwords. In the end, though, these details do not matter so much in our implementation.

Step 2: Review the naïve Bayes Classifier

To classify the email messages, we use a multinomial Naïve Bayes model (where the word frequency is used in feature vectors, instead of a Boolean variable that indicates the word occurs or not). Review the slides, including the smoothing method and the evaluation measures.

Step 3: Implementing Naïve Bayes

Feature Selection: Here are some guidelines that will help you to generate your own features.

Data contents

The data pack you downloaded contains 4 folders.

a. The folders "nonspam-train" and "spam-train" contain the preprocessed emails you will use for training. They each have 350 emails.

b. The folders "nonspam-train" and "nonspam-test" constitute the test set containing 130 spam and 130 nonspam emails. These are the documents you will make predictions on. Notice that even though the separate folders tell you the correct labeling, you should make your predictions on all the test documents without this knowledge. After you make your predictions, you can use the correct labeling to check whether your classifications were correct.

Dictionary/Vocabulary

You will need to generate a dictionary for your model. There is more than one way to do this, but an easy method is to count the occurrences of all words that appear in the emails and choose your dictionary to be the most frequent words. You are suggested to pick the dictionary to be the **2,500** most frequent words. Remember to take the counts over all of the emails: spam and nonspam in the **training set**.

Once you have the dictionary, you will need to represent your documents as feature vectors over the space of the dictionary words. One possible way to do so:

For each document, keep track of the dictionary words that appear, along with the count of the number of occurrences.

Training and testing

Finally, you will need to train your model on the training set and predict the spam/nonspam classification on the test set. **You are suggested to have separate scripts for training and testing.** That way, after you trained your model (that is, after calculated all the parameters of the model), you can run the testing independently as long as the parameters are stored.

Calculate the parameters (the necessary probabilities) from the training data.

Using the model parameters you obtained from training, classify each test document as spam or non-spam.

Note: Be sure you work with log probabilities.

Evaluation

Draw the 2x2 table to report the number of false positive, true positives, false negatives, and true negatives.

Report the precision, recall and f1-score of your system.

[bonus question: 25pts] **Smaller Training sets:**

Let's see how the classification performance changes when you train on smaller training sets, but test on the same test set as before. So far you have been working with a 700-document training set. You will now modify your program to train on 50, 100, and 400 documents (the spam to nonspam ratio will still be one-to-one).

Select email subsets of 50, 100, and 400, keeping each subset 50% spam and 50% nonspam. For each of these subsets, generate the training features as you did before and train your model. Then, test your model on the 260-document test set and record your classification performance.